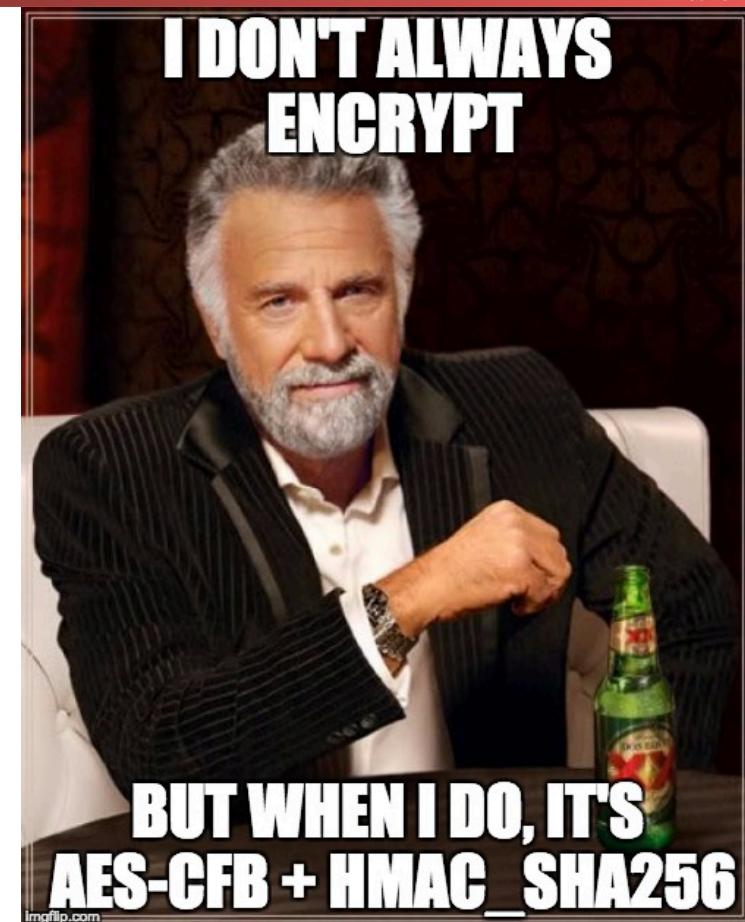
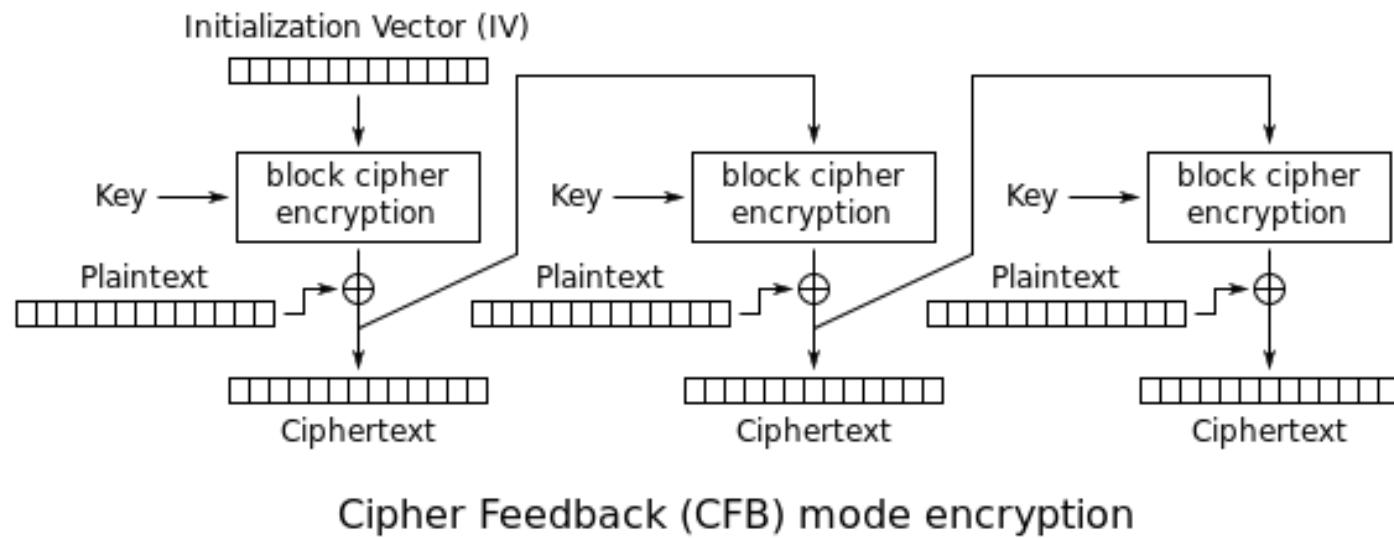


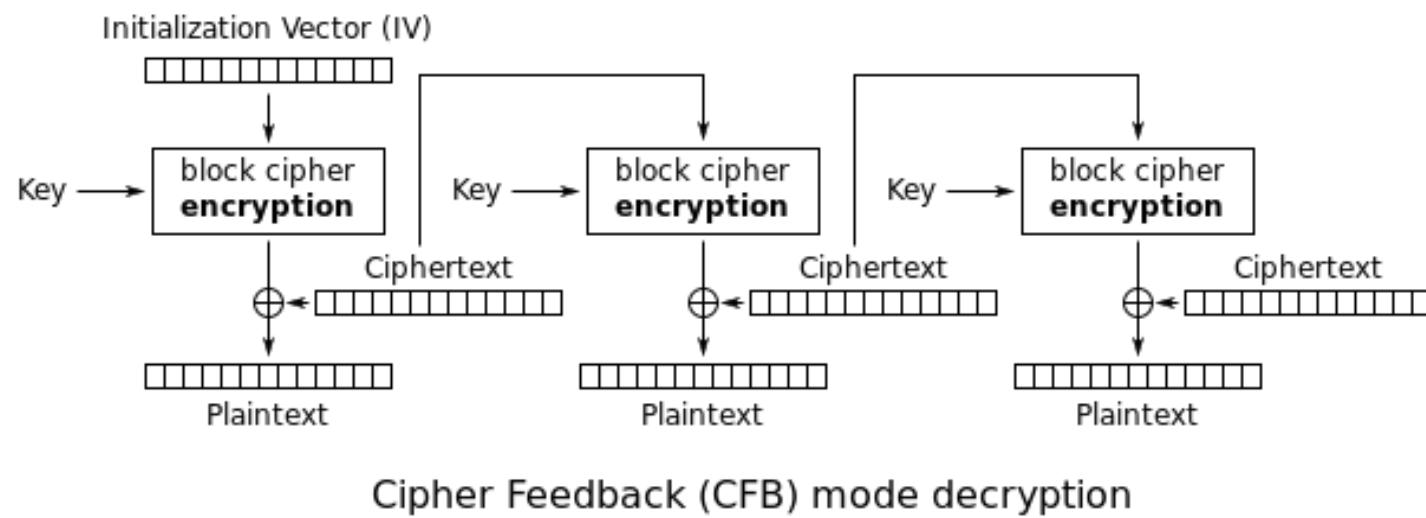
Crypto 2



CFB Encryption



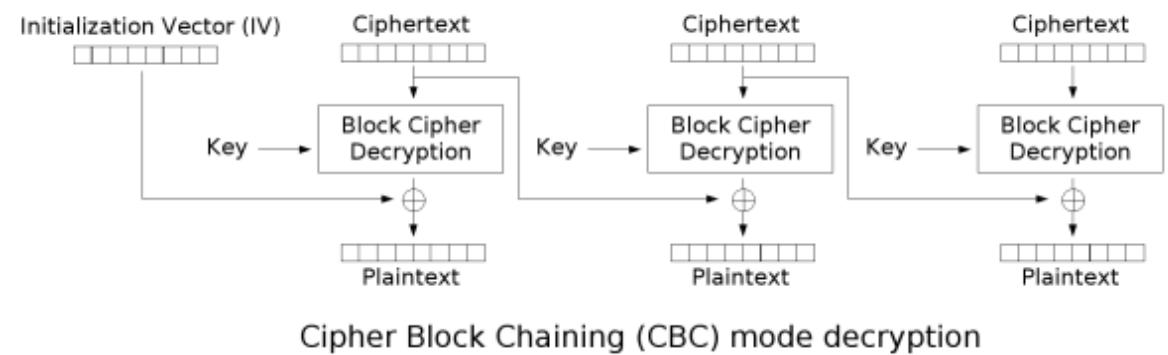
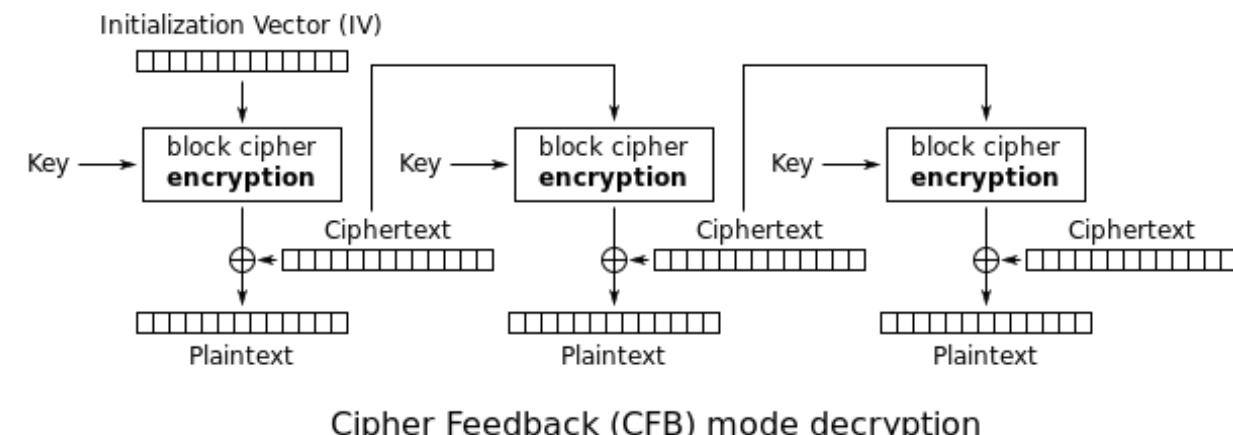
CFB Decryption



CFB doesn't need to pad...

- Since the encryption is XORed with the plaintext...
 - You can end on a "short" block without a problem
 - So more convenient than CBC mode
- But similar security properties as CBC mode
 - Sequential encryption, parallel decryption
 - Same error propagation effects
 - Effectively the same for IND-CPA
- But a bit worse if you reuse the IV

Error Propagation



Mallory the Manipulator

- Mallory is an active attacker
 - Can introduce new messages (ciphertext)
 - Can “replay” previous ciphertexts
 - Can cause messages to be reordered or discarded
- A “Man in the Middle” (MITM) attacker
 - Can be much more powerful than just eavesdropping



Encryption Does Not Provide Integrity

- Simple example: Consider a block cipher in CTR mode...
- Suppose Mallory knows that Alice sends to Bob “Pay Mal \$0100”. Mallory intercepts corresponding C
 - $M = \text{"Pay Mal $0100"}$. $C = \text{"r4ZC#jj8qThMK"}$
 - $M_{10..13} = \text{"0100"}$. $C_{10..13} = \text{"ThMK"}$
 - Mallory wants to replace some bits of C...



Encryption Does Not Provide Integrity

- Mallory computes
 - “0100” \oplus $C_{10..13}$
 - Tells Mallory that section of the counter XOR:
Remember that CTR mode computes $E_k(IV||CTR)$ and XORs it with the corresponding part of the message
 - $C'_{10..13} = "9999" \oplus "0100" \oplus C_{10..13}$
 - Mallory now forwards to Bob a full $C' = C_{0..9}||C'_{10..13}||C_{14...}$
 - Bob will decrypt the message as "Pay Mal \$9999"...
 - For a CTR mode cipher, Mallory can in general replace any *known* message M with a message M' of equal length!

Integrity and Authentication

- Integrity: Bob can confirm that what he's received is exactly the message M that was originally sent
- Authentication: Bob can confirm that what he's received was indeed generated by Alice
- Reminder: for either, confidentiality may-or-may-not matter
 - E.g. conf. not needed when Mozilla distributes a new Firefox binary
- Approach using symmetric-key cryptography:
 - Integrity via MACs (which use a shared secret key K)
 - Authentication arises due to confidence that only Alice & Bob have K
- Approach using public-key cryptography (later on):
 - “Digital signatures” provide both integrity & authentication together
- Key building block: cryptographically strong hash functions

Hash Functions

- Properties
 - Variable input size
 - Fixed output size (e.g., 256 bits)
 - Efficient to compute
 - Pseudo-random (mixes up input extremely well):
A single bit changes on the input and ~1/2 the bits should change on the output
- Provides a “fingerprint” of a document
 - E.g. “`shasum -a 256 <exams/mt1-solutions.pdf`” prints
`0843b3802601c848f73ccb5013afa2d5c4d424a6ef477890ebf8db9bc4f7d13d`

Cryptographically Strong Hash Functions

- A collision occurs if $x \neq y$ but $\text{Hash}(x) = \text{Hash}(y)$
 - Since input size > output size, collisions do happen
- A cryptographically strong $\text{Hash}(x)$ provides three properties:
 - One-way: $h = \text{Hash}(x)$ easy to compute, but not to invert.
 - Intractable to find *any* x' s.t. $\text{Hash}(x') = h$, for a given h
 - Also termed “preimage resistant”

$H(\text{🐮}) =$



Cryptographically Strong Hash Functions

- The other two properties of a cryptographically strong **Hash(x)**:
 - Second preimage resistant: given x , intractable to find x' s.t. $\text{Hash}(x) = \text{Hash}(x')$
 - Collision resistant: intractable to find any x, y s.t. $\text{Hash}(x) = \text{Hash}(y)$
- Collision resistant \implies Second preimage resistant
 - We consider them separately because given Hash might differ in how well it resists each
 - Also, the Birthday Paradox means that for n-bit Hash, finding $x-y$ pair takes only $\approx 2^{n/2}$ hashes
 - Vs. potentially 2^n tries for x' : $\text{Hash}(x) = \text{Hash}(x')$ for given x
- Plus a hash function should look "random"
 - A "PRF" or Pseudo-Random Function

Cryptographically Strong Hash Functions, con't

- Some contemporary hash functions
 - MD5: 128 bits
 - broken – lack of collision resistance
 - Collisions for the heck of it: <https://shells.aachen.ccc.de/~spq/md5.gif>
An MD5 "hash quine": an animated GIF that shows its own hash
 - SHA-1: 160 bits broken spring 2017, but was known to be weak yet still used...
 - SHA-256/SHA-384/SHA-512: 256, 384, 512 bits in the SHA-2 family, at least not currently broken
 - SHA-3: New standard! Yayyy!!!! (Based on Keccak, again 256b, 384b, and 512b options)
- Provide a handy way to unambiguously refer to large documents
 - If hash can be securely communicated, provides integrity
 - E.g. Mozilla securely publishes SHA-256(new FF binary)
 - Anyone who fetches binary can use "`cat binary | shasum -a 256`" to confirm it's the right one, untampered
- Not enough by themselves for integrity, since functions are completely known –
Mallory can just compute revised hash value to go with altered message

SHA-256...

- SHA-256/SHA-384 are two parameters for the SHA-2 hash algorithm, returning 256b or 384b hashes
- Works on blocks with a truncation routine to make it act on sequences of arbitrary length
- Is vulnerable to a ***length-extension attack***: s is secret
 - Mallory knows $\text{len}(s)$, $H(s)$
 - Mallory can use this to calculate $H(s||M)$ for an M of Mallory's construction
 - Works because ***all the internal state*** at the point of calculating $H(s||...)$ is derivable from $H(s)$ and $\text{len}(s)$
 - New SHA-3 standard (Keccak) does not have this property

Stupid Hash Tricks: Sample A File...

- BlackHat Dude claims to have 150M records stolen from Equifax...
 - How can I as a reporter verify this?
 - Idea: If I can have the hacker select 10 *random* lines...
 - And in selecting them also say something about the size of the file...
 - Voila! Verify those lines and I now know he's not full of BS
 - Can I use hashing to write a small script which the BlackHat Dude can run?
 - Where I can easily verify that the 10 lines were sampled at random, and can't be faked?

Sample a File

```
#!/usr/bin/env python
import hashlib, sys
hashes = {}

for line in sys.stdin:
    line = line.strip()
    for x in range(10):
        tmp = "%s-%i-nickrocks" % (line, x)
        hashval = hashlib.sha256(tmp)
        h = hashval.digest()
        if x not in hashes or hashes[x][0] > h:
            hashes[x] = (h, hashval, tmp)

for x in range(10):
    h, hashval, val = hashes[x]
    print "%s=%\"%s\\"" % (hashval.hexdigest(), val)
```

Why does this work?

- For each x in range 0-9...
 - Calculates $H(\text{line}||x)$
 - Stores the lowest hash matching so far
- Since the hash appears random...
 - Each iteration is an *independent* selection from the file
 - The expected value of $H(\text{line}||x)$ is a function of the size of the file:
More lines, and the value is smaller
- To fake it...
 - Would need to generate fake lines, *and see if the hash is suitably low*
 - Yet would need to make sure these fake lines semantically match!
 - Thus you can't just go "John Q Fake", "John Q Fakke", "Fake, John Q", etc...

Message Authentication Codes (MACs)

- Symmetric-key approach for integrity
 - Uses a shared (secret) key K
- Goal: when Bob receives a message, can confidently determine it hasn't been altered
 - In addition, whomever sent it must have possessed K
(\Rightarrow message authentication, sorta...)
- Conceptual approach:
 - Alice sends $\{M, T\}$ to Bob, with tag $T = \text{MAC}(K, M)$
 - Note, M could instead be $C = E_K(M)$, but not required
 - When Bob receives $\{M', T'\}$, Bob checks whether $T' = \text{MAC}(K, M')$
 - If so, Bob concludes message untampered, came from Alice
 - If not, Bob discards message as tampered/corrupted

Requirements for Secure MAC Functions

- Suppose MITM attacker Mallory intercepts Alice's $\{M, T\}$ transmission ...
 - ... and wants to replace M with altered M^*
 - ... but doesn't know shared secret key K
- We have secure integrity if MAC function $T = \text{MAC}(M, K)$ has two properties:
 - Mallory can't compute $T^* = \text{MAC}(M^*, K)$
 - Otherwise, could send Bob $\{M^*, T^*\}$ and fool him
 - Mallory can't find M^{**} such that $\text{MAC}(M^{**}, K) = T$
 - Otherwise, could send Bob $\{M^{**}, T\}$ and fool him
- These need to hold even if Mallory can observe many $\{M_i, T_i\}$ pairs, including for M_i 's she chose

MAC then Encrypt or Encrypt then MAC

- You should ***never*** use the same key for the MAC and the Encryption
 - Some MACs will break completely if you reuse the key
 - Even if it is ***probably*** safe (eg, AES for encryption, HMAC for MAC) its still a bad idea
- MAC then Encrypt:
 - Compute $T = \text{MAC}(M, K_{\text{mac}})$, send $C = E(M||T, K_{\text{encrypt}})$
- Encrypt then MAC:
 - Compute $C = E(M, K_{\text{encrypt}})$, $T = \text{MAC}(M, K_{\text{mac}})$, send $C||T$
- Theoretically they are the same, but...
 - Once again, its time for ...



HTTPS Authentication in Practice

- When you log into a web site, it sets a "cookie" in your browser
 - All subsequent requests include this cookie so the web server knows who you are
- If an attacker can get your cookie...
 - They can impersonate you on the "Secure" site
- And the attacker can create multiple tries
 - On a WiFi network, inject a bit of JavaScript that repeatedly connects to the site
 - While as a man-in-the-middle to manipulate connections



The TLS 1.0 "Lucky13" Attack: "F-U, This is Cryptography"

- HTTPS/TLS uses MAC then Encrypt
 - With CBC encryption
- The Lucky13 attack changes the cipher text in an attempt to discover the state of a byte
 - But can't predict the MAC
 - The TLS connection retries after each failure so the attacker can try multiple times
 - Goal is to determine the status each byte in the authentication cookie which is in a known position
- It detects the ***timing*** of the error response
 - Which is different if the guess is right or wrong
 - Even though the underlying algorithm was "***proved***" secure!
- So always do Encrypt then MAC since, once again, it is more mistake tolerant



The best MAC construction: HMAC

- Idea is to turn a hash function into a MAC
 - Since hash functions are often much faster than encryption
 - While still maintaining the properties of being a cryptographic hash
- Reduce/expand the key to a single hash block
- XOR the key with the i_pad
 - 0x363636... (one hash block long)
- Hash ((K \oplus i_pad) || message)
- XOR the key with the o_pad
 - 0x5c5c5c...
- Hash ((K \oplus o_pad) || first hash)

```
function hmac (key, message) {
    if (length(key) > blocksize) {
        key = hash(key)
    }
    while (length(key) < blocksize) {
        key = key || 0x00
    }
    o_key_pad = 0x5c5c...  $\oplus$  key
    i_key_pad = 0x3636...  $\oplus$  key
    return hash(o_key_pad ||
                hash(i_key_pad || message))
}
```

Why This Structure?

- `i_pad` and `o_pad` are slightly arbitrary
 - But it is necessary for security for the two values to be different
 - So for paranoia chose very different bit patterns
- Second hash prevents appending data
 - Otherwise attacker could add more to the message and the HMAC and it would still be a valid HMAC for the key
 - Wouldn't be a problem with the key at the **end** but at the start makes it easier to capture intermediate HMACs
- Is a Pseudo Random Function if the underlying hash is a PRF
 - AKA if you can break this, you can break the hash!

```
function hmac (key, message) {  
    if (length(key) > blocksize) {  
        key = hash(key)  
    }  
    while (length(key) < blocksize) {  
        key = key || 0x00  
    }  
    o_key_pad = 0x5c5c... ⊕ key  
    i_key_pad = 0x3636... ⊕ key  
    return hash(o_key_pad ||  
               hash(i_key_pad || message))  
}
```

Great Properties of HMAC...

- It is still a hash function!
 - So all the good things of a cryptographic hash:
Nobody should be able to calculate **M** given **HMAC(M,K)** if they know **K** but don't know **M**
 - An attacker who doesn't know **K** *can't even verify* if **HMAC(M,K) == M**
 - Very different from the hash alone, and potentially very useful:
Attacker can't even brute force try to find **M** based on **HMAC(M,K)** if they don't know **K**
- Its probably safe if you screw up and use the same key for both MAC and Encrypt
 - Since it is a different algorithm than the encryption function...
 - ***But you shouldn't do this anyway!***

Considerations when using MACs

- Along with messages, can use for data at rest
 - E.g. laptop left in hotel, providing you don't store the key on the laptop
 - Can build an efficient data structure for this that doesn't require re-MAC'ing over entire disk image when just a few files change
- MACs in general provide no promise not to leak info about message
 - Compute MAC on ciphertext if this matters
 - Or just use HMAC, which **does** promise not to leak info if the underlying hash function doesn't
- **NEVER** use the same key for MAC and Encryption...
 - Known "FU-this-is-crypto" scenarios reusing an encryption key for MAC in some algorithms when its the same underlying block cipher for both

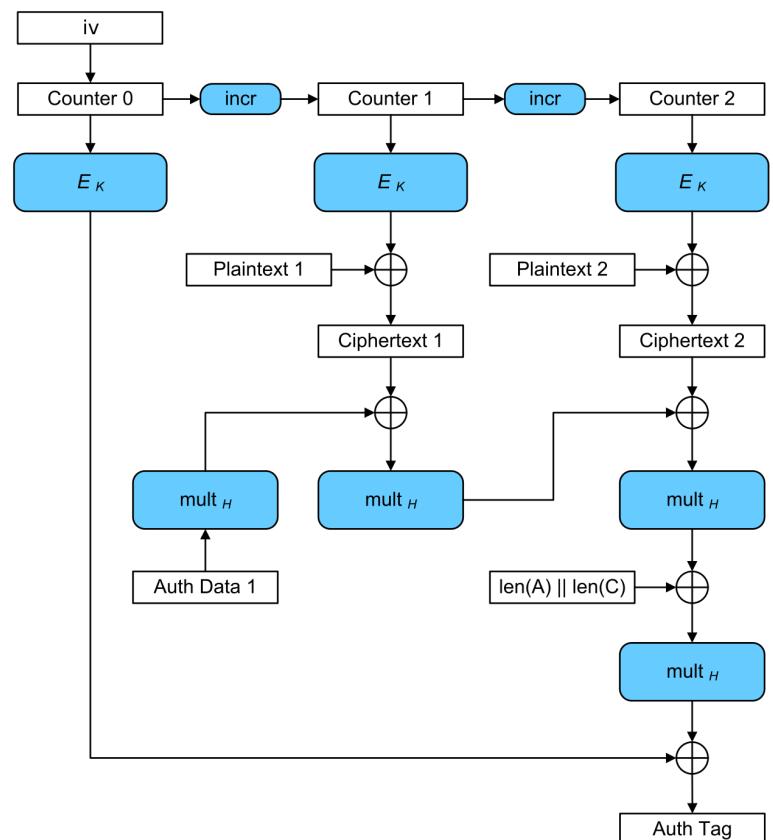


Plus AEAD Encryption Modes...

- The latest block cipher modes are "AEAD":
 - Authenticated Encryption with Additional Data
- Provides both **integrity** **and** confidentiality over the data
 - With **integrity** also provided for the "Additional Data"
- Used right, these are great
 - Assuming you use a library...
- Used wrong...
 - The AEAD modes are built for "performance", which means parallelization, which means CTR mode, which means IV reuse is a disaster!

GCM (Galios Counter Mode) AEAD Counter Mode

- E_k is just standard encryption
- mult_h is 128b multiplication over a special field
 - Not going into the details of the magic math because I don't understand it myself
- VERY fast mode of operation
 - Fully parallel encryption
 - Galios multiplication isn't parallelizeable but it is very light weight
 - Far lighter weight than HMAC-SHA256
- But a huge pitfall....
 - If you reuse the IV, not only do you lose confidentiality (like CTR mode)...
You lose integrity as well!
 - Which wouldn't happen if you used CTR-HMAC-SHA256



A Lot Of Uses for Random Numbers...

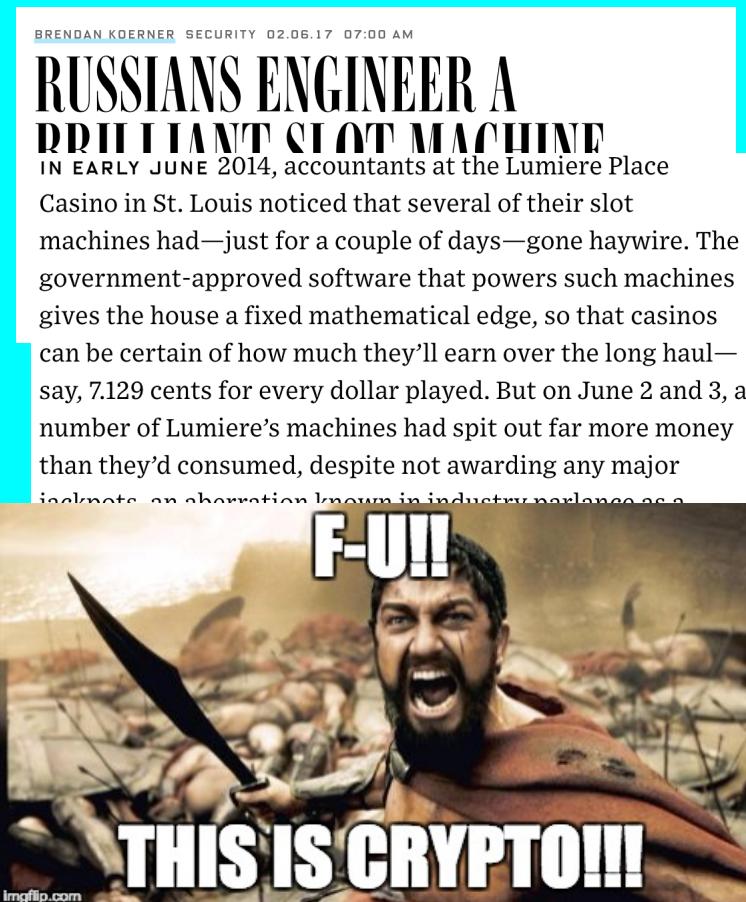
- The key foundation for all modern cryptographic systems is often not encryption but these "random" numbers!
- So many times you need to get something random:
 - A random cryptographic key
 - A random initialization vector
 - A "nonce" (use-once item)
 - A unique identifier
 - Stream Ciphers
- If an attacker can *predict* a random number things can catastrophically fail

Breaking Slot Machines

Computer Science 161

Weaver

- Some casinos experienced unusual bad "luck"
 - The suspicious players would wait and then all of a sudden try to play
- The slot machines have ***predictable*** pRNG
 - Which was based on the current time & a seed
- So play a little...
 - With a cellphone watching
 - And now you know when to press "spin" to be more likely to win
- Oh, and this ***never*** effected Vegas!
 - ***Evaluation standards*** for Nevada slot machines specifically designed to address this sort of issue



Breaking Bitcoin Wallets

Computer Science 161

Weaver

- blockchain.info supports "web wallets"
 - Javascript that protects your Bitcoin
- The private key for Bitcoin needs to be random
 - Because otherwise an attacker can spend the money
- An "Improvement" [sic] to the RNG reduced the entropy (the actual randomness)
 - Any wallet created with this improvement was brute-forceable and could be stolen

Improvements to RNG

zootreeves committed on Dec 7, 2014

1 parent b0d5639

Showing 1 changed file with 26 additions and 28 deletions.

S4 bitcoinjs-lib/src/jspb/rng.js

```
@@ -8,15 +8,16 @@ var rng_state;
 8   8   var rng_pool;
 9   9   var rng_pptr;
10  10
11 // Mix in a 32-bit integer into the pool
12 -function rng_seed_int(x) {
13 -  rng_pool[rng_pptr++] ^= x & 255;
14 -  rng_pool[rng_pptr++] ^= (x >> 8) & 255;
15 -  rng_pool[rng_pptr++] ^= (x >> 16) & 255;
16 -  rng_pool[rng_pptr++] ^= (x >> 24) & 255;
```



TRUE Random Numbers

- True random numbers generally require a physical process
- Common circuit is an unusable ring oscillator built into the CPU
 - It is then sampled at a low rate to generate true random bits which are then fed into a pRNG on the CPU
- Other common sources are human activity measured at very fine time scales
 - Keystroke timing, mouse movements, etc
 - "Wiggle the mouse to generate entropy for a key"
 - Network/disk activity which is often human driven
- More exotic ones are possible:
 - Cloudflare has a wall of lava lamps that are recorded by a HD video camera which views the lamps through a rotating prism: It is just one source of the randomness



Combining Entropy

- Many physical entropy sources are biased
 - Some have significant biases: e.g. a coin that flips "heads" 90% of the time!
 - Some aren't very good: e.g. keystroke timing at a microsecond granularity
- The general procedure is to combine various sources of entropy
- The goal is to be able to take multiple crappy sources of entropy
 - Measured in how many bits:
A single flip of a fair coin is 1 bit of entropy
 - And combine into a value where the entropy is the minimum of the sum of all entropy sources (maxed out by the # of bits in the hash function itself)
 - **N-1** bad sources and **1** good source -> good pRNG state

Pseudo Random Number Generators (aka Deterministic Random Bit Generators)

- Unfortunately one needs a *lot* of random numbers in cryptography
 - More than one can generally get by just using the physical entropy source
- Enter the pRNG or DRBG
 - If one knows the state it is entirely predictable
 - If one doesn't know the state it should be indistinguishable from a random string
- Three operations
 - Instantiate: (aka Seed) Set the internal state based on the real entropy sources
 - Reseed: Update the internal state based on both the previous state and **additional entropy**
 - The big different from a simple stream cipher
 - Generate: Generate a series of random bits based on the internal state
 - Generate can also optionally add in additional entropy
- **instantiate(entropy)**
reseed(entropy)
generate(bits, {optional entropy})

Properties for the pRNG

- Can a pRNG be truly random?
 - No. For seed length s , it can only generate at most 2^s distinct possible sequences.
- A cryptographically strong pRNG “looks” truly random to an attacker
 - Attacker ***cannot distinguish*** it from a random sequence:
If the attacker can tell a sufficiently long bitstream was generated by the pRNG instead of a truly random source it isn't a good pRNG

Prediction and Rollback Resistance

- A pRNG should be predictable only if you know the internal state
 - It is this predictability which is why its called "pseudo"
- If the attacker does not know the internal state
 - The attacker should not be able to distinguish a truly random string from one generated by the pRNG
- It **should** also be rollback-resistant
 - Even if the attacker finds out the state at time T, they should not be able to determine what the state was at T-1
 - More precisely, if presented with two random strings, one truly random and one generated by the pRNG at time T-1, the attacker should not be able to distinguish between the two
 - Rollback resistance isn't specifically required in a pRNG...
But it should be

Why "Rollback Resistance" is Essential

- Assume attacker, at time T , is able to obtain all the internal state of the pRNG
 - How? E.g. the pRNG screwed up and instead of an IV, released the internal state, or the pRNG is bad...
 - Attacker observes how the pRNG was used
 - T_{-1} = Random Session key
 T_0 = Nonce/IV
 - Now if the pRNG doesn't resist rollback, and the attacker gets the state at T_0 , attacker can know the session key! And we are back to...



More on Seeding and Reseeding

- Seeding should take all the different physical entropy sources available
 - If one source has 0 entropy, it ***must not*** reduce the entropy of the seed
 - We can shove a whole bunch of low-entropy sources together and create a high-entropy seed
- Reseeding ***adds*** in even more entropy
 - **F(internal_state, new material)**
 - Again, even if reseeding with 0 entropy, it ***must not*** reduce the entropy of the seed

Probably the best pRNG/DRBG: HMAC_DRBG

- Generally believed to be the best
 - ***Accept no substitutes!***
- Two internal state registers, V and K
 - Each the same size as the hash function's output
- V is used as (part of) the data input into HMAC, while K is the key
- If you can break this pRNG you can ***either break the underlying hash function or break a significant assumption about how HMAC works***
 - Yes, security proofs sometimes are a very good thing and actually do work
 - So as long as the security proof for HMAC is correct, the security proof for HMAC_DRBG is correct!

HMAC_DRBG

Generate

- The basic generation function
- Remarks:
 - It requires one HMAC call per blocksize-bits of state
 - Then two more HMAC calls to update the internal state
- Prediction resistance:
 - If you can distinguish new **K** from random when you don't know old **K**:
You've distinguished HMAC from a random function!
Which means you've either broken the hash or the HMAC construction
- Rollback resistance:
 - If you can learn old **K** from new **K** and **V**:
You've reversed the hash function!

```
function hmac_drbg_generate (state, n) {  
    tmp = ""  
    while(len(tmp) < N) {  
        state.v = hmac(state.k, state.v)  
        tmp = tmp || state.v  
    }  
    // Update state with no input  
    state.k = hmac(state.k, state.v || 0x00)  
    state.v = hmac(state.k, state.v)  
    // Return the first N bits of tmp  
    return tmp[0:N]  
}
```

HMAC_DRBG

Update

- Used instead of the "no-input update" when you have additional entropy on the generate call
- Used standalone for both instantiate (**state.k = state.v = 0**) and reseed (keep **state.k** and **state.v**)
- Designed so that even if the attacker controls the input but doesn't know **k**:
The attacker should not be able to predict the new **k**

```
function hmac_drbg_update (state, input) {  
    state.k = hmac(state.k, state.v || 0x00  
                  || input)  
    state.v = hmac(state.k, state.v)  
    state.k = hmac(state.k, state.v || 0x01  
                  || input)  
    state.v = hmac(state.k, state.v)  
}
```

Generating ***true*** random numbers

- Modern CPUs have true random number generators
 - Sample a noisy circuit at a low rate or similar tricks
- These sources are biased...
 - They are also ***slow***
- So use this as an entropy source to feed a pRNG on the chip
 - Now you can get random numbers quickly
- Very fast
- Vulnerable to tampering!
 - You can't actually test that the pRNG circuit is 100% correct without adding paths that could potentially sabotage the pRNG circuit
 - Sabotage that can reduce effective entropy to 32b are possible